

ASSIGNMENT 4

Q1

Errors in each technique:

```
For sklearn LinearRegression :  
  
RMSE:  0.6207498618060853  
MAE:   0.5126346094502968  
-----  
For linear regression using normal equations :  
  
RMSE:  0.6207498618060854  
MAE:   0.5126346094502968  
-----  
For linear regression using SVD :  
  
RMSE:  0.6207498618060853  
MAE:   0.5126346094502968  
-----
```

Time for each technique :

```
Time taken by Sklearn's implementation: 0.0027539730072021484 seconds  
Time taken by normal equations: 0.0002944469451904297 seconds  
Time taken by SVD: 0.0048351287841796875 seconds
```

Q2

```
\ / \
Batch Gradient Descent with manual gradient computation for unregularized objective :
Batch size= 1 , RMSE: 0.7905046278690018
Batch size= 1 , MAE: 0.6314635711163878
-----
Ridge regression with manual gradient computation :
Batch size= 1 , RMSE: 0.8254707844072325
Batch size= 1 , MAE: 0.6304058871895001
-----
Batch gradient descent with jax gradient computation for unregularized objective :
Batch size= 1 , RMSE: 0.7995486600057372
Batch size= 1 , MAE: 0.6487486828744025
-----
Ridge regression with jax gradient computation :
Batch size= 1 , RMSE: 0.7917161721091319
Batch size= 1 , MAE: 0.6298493002999888
-----
LASSO :
Batch size= 1 , RMSE: 0.832308215077365
Batch size= 1 , MAE: 0.6957399363046195
-----
SGD with momentum for manual gradient computation on unregularized objective :
Batch size= 1 , RMSE: 0.8055490409646227
Batch size= 1 , MAE: 0.6268142945465978
-----
Ridge regression on SGD with momentum using manual computation :
Batch size= 1 , RMSE: 0.9785720532637601
Batch size= 1 , MAE: 0.8140122430647128
-----
SGD with momentum for jax gradient on unregularized:
Batch size= 1 , RMSE: 0.8892927066451223
Batch size= 1 , MAE: 0.7061438121083694
-----
```

SGD with momentum for jax gradient on unregularized:

Batch size= 1 , RMSE: 0.8892927066451223

Batch size= 1 , MAE: 0.7061438121083694

Ridge regression using jax using SGD :

Batch size= 1 , RMSE: 0.7946100588309402

Batch size= 1 , MAE: 0.6417867355312361

LASSO using SGD:

Batch size= 1 , RMSE: 0.8280695290241143

Batch size= 1 , MAE: 0.6708981730103792

Final data:

	Regressor	Method of calculating gradient	Penalty type \
0	Batch gradient descent	manual	unregularized
1	Batch gradient descent	manual	unregularized
2	Batch gradient descent	manual	unregularized
3	Batch gradient descent	manual	unregularized
4	Batch gradient descent	manual	unregularized
..
395	SGD with momentum	jax	l1
396	SGD with momentum	jax	l1
397	SGD with momentum	jax	l1
398	SGD with momentum	jax	l1
399	SGD with momentum	jax	l1

	Learning rate	Lambda	Momentum	RMSE	MAE	Training time
0	0.001	0.001	NA	0.869777	0.681868	0.169184
1	0.001	0.010	NA	1.204946	1.018111	0.170525
2	0.001	0.100	NA	1.497003	1.245115	0.158658
3	0.001	1.000	NA	1.508113	1.172299	0.177218
4	0.010	0.001	NA	0.805932	0.648113	0.168406

396	1.000	1.000	-3	NaN	NaN	1.438138
397	1.000	1.000	-2	NaN	NaN	1.431901
398	1.000	1.000	-1	NaN	NaN	1.490301
399	1.000	1.000	0	NaN	NaN	1.710637

[400 rows x 9 columns]

Optimal hyperparameters for batch gradient descent based on RMSE:

	Regressor	Method of calculating gradient	Penalty type	
58	Batch gradient descent	jax	l2	

	Learning rate	Lambda	Momentum	RMSE	MAE	Training time
58	0.1	0.1	NA	0.792424	0.618377	2.312205

Optimal Hyperparameters for batch gradient descent based on MAE

	Regressor	Method of calculating gradient	Penalty type	
58	Batch gradient descent	jax	l2	

	Learning rate	Lambda	Momentum	RMSE	MAE	Training time
58	0.1	0.1	NA	0.792424	0.618377	2.312205

Optimal hyperparameters for SGD with momentum based on RMSE

	Regressor	Method of calculating gradient	Penalty type	
296	SGD with momentum	jax	l2	


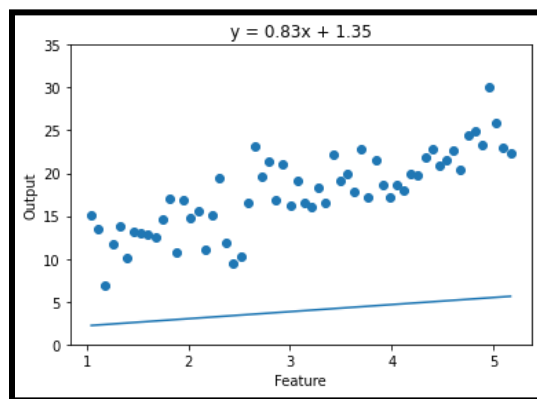


	Learning rate	Lambda	Momentum	RMSE	MAE	Training time
296	0.01	0.1	-3	0.790582	0.62824	1.553014

Optimal hyperparameters for SGD with momentum based on MAE

	Regressor	Method of calculating gradient	Penalty type	
250	SGD with momentum	jax	unregularized	


	Learning rate	Lambda	Momentum	RMSE	MAE	Training time
250	0.1	0.1	-1	0.800185	0.617581	0.979765

Q3 Gif uploaded on Git

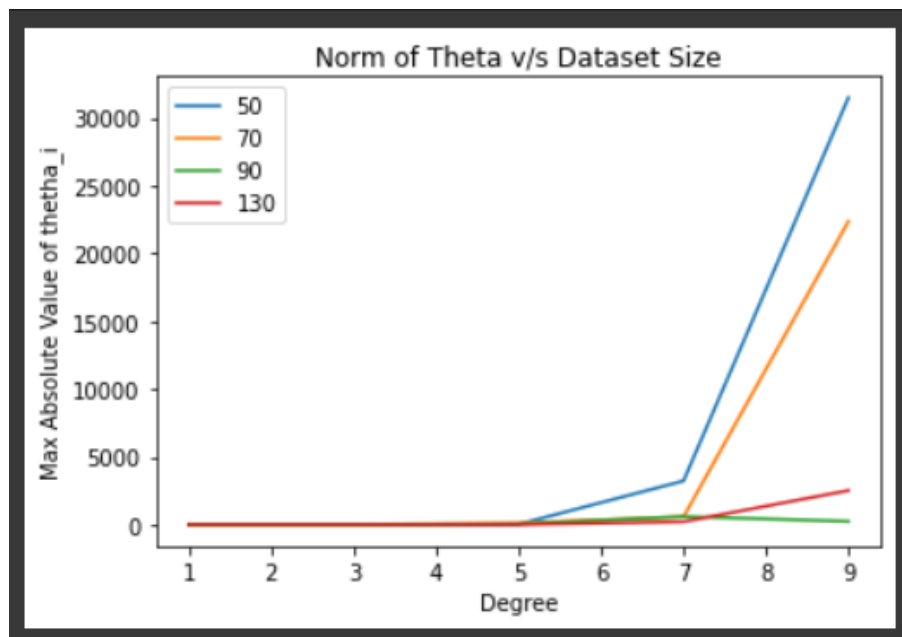


Q4

Norm vs Degree

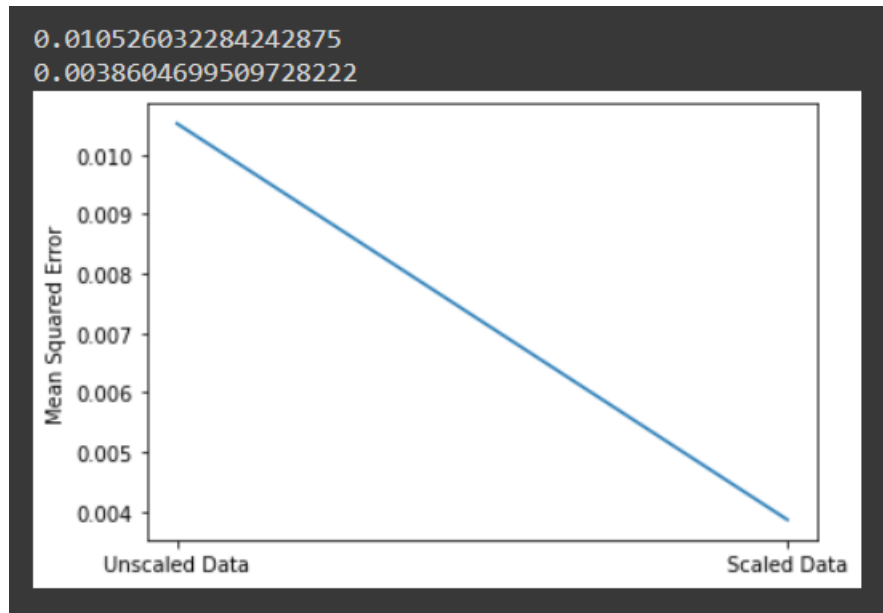


Q5




Q6

For random dataset:



For diabetes dataset:



Q7

